

Consentir: An Electronic Patient Consent Management System

Atif Khan, Sarah Nadi
David R. Cheriton School of Computer Science
University of Waterloo
Ontario, Canada
a78khan, snadi@uwaterloo.ca

Abstract—Effective management of patient information plays a critical role in providing appropriate health care. As Information Communication Technologies (ICTs) are being integrated into the medical domain, patient information is becoming increasingly managed through electronic systems. This poses the challenge of ensuring patient privacy in accordance with each patient's consent policy. Therefore, there is a need for an electronic patient consent management system which can capture patient consent, process it, and further decide who can access a patient's information. In this paper, we propose such a system, *Consentir*, where patient information and consent policies, as well as other access policies are represented in Notation 3 (N3). A reasoning engine, Euler, is used to determine if access should be granted or denied, according to the patient's consent policy. We built a prototype to show how semantic web based techniques can be used to solve the issue of patient consent. We currently support five different consent policies: opt in, opt in with exceptions to specific people, opt in with exceptions to sensitive information, opt out, and opt out with emergency override. In our prototype, we demonstrate how these consent policies affect information access in different health care situations through twelve different scenarios.

I. INTRODUCTION

With the advance of Information Communication Technologies (ICTs), more medical organizations are utilizing electronic systems to capture, manage and use patient related information. The scope of this information varies from organization wide Electronic Medical Records (EMRs) to Electronic Health Records (EHRs) shared between different organizations. Although this improves the effectiveness of information exchange, coordination, and use, it also raises the critical issue of patient privacy. Usually, patient information is collected for the primary purpose of providing health care for a specific episode. Any secondary use must be in accordance with patient consent. It has been argued that a patient should be aware of all the systems collecting their information, and should be able to specify how this information can be used [1]. Accordingly, each patient should choose a consent policy that reflects how they would like their medical information to be used. Ideally, in an electronic information system, the system would automatically grant or deny permission to access a patient's record according to the corresponding consent policy. However, it is often difficult, if not impossible, to predict all future-use scenarios and enforce patient consent in an appropriate manner. It should be noted that in this context,

patient consent is related to accessing medical information. This is different from the patient consent required to provide certain treatments or perform specific procedures.

There are many aspects of this problem that need to be solved. First, how can patient consent be effectively and accurately captured electronically? Second, how can the captured consent policies be represented and processed internally? Finally, how can we define consent policies that protect a patient's privacy, but do not compromise the individual's health at the same time? All three problems are very important in this domain. However, in this paper, we focus on the second problem, and briefly touch upon the third.

To address these problems, we built *Consentir*, a policy (rule) based patient consent management system. This system utilizes (previously captured) patient consent information and various operational policies as input. Each policy will be represented by a set of RDF rules in Notation 3 (N3) [2]. We will use the Euler proof mechanism [3] to compute (a) the result and (b) the proof of the aggregated rules.

Our goal is to demonstrate that access to patient information can be protected in a real-time manner by utilizing policy based consent management. The advantage of using Euler to generate a proof is that in the future, proofs can be validated between different systems when exchanging EHRs. This offers a major improvement over the current industry practice of patient consent management and secondary use of patient information. This paper describes the preliminary effort in this direction where we apply our idea to a selection of consent policies and situations, and produce a working prototype.

The contributions of this work are as follows:

- A set of consent policies represented in N3 notation. These are expressed as N3 rules which allow or deny access to the specified documents.
- A collection of executable scenarios that show how the consent policies are applied in different situations.
- A working prototype, implemented in Java, to demonstrate these scenarios.

The rest of this paper is organized as follows. Section II provides background information about patient consent, N3 notation, and the Euler engine. Section III describes some of the related work that has been done to develop an electronic patient consent system. Section IV explains the different consent policies, and specifies those included in our prototype.

Section V describes *Consentir*, the prototype system developed in this work. It gives an overview of the system as well as the policy rules and facts used. Section VI discusses some of the current limitations of our system and possible future work that can be done to address them. Finally, Section VII concludes this paper.

II. BACKGROUND

A. Patient Consent

The Consent Challenge: Effective patient management involves collaboration between multiple parties, each making its contribution to the overall health care provided to the patient. An effective health care framework will not only support multiple (diverse) participants, but it will also ensure proper information exchange between these participants to provide the highest level of service (care). Although the prevalent exchange of patient information drastically improves the care provided, it does raise a major challenge of information management. More specifically, patients would like to maintain control over the various attributes of their private and confidential personal information. The attributes being access, collaboration, and revocation. We refer to this as the “Consent Challenge” in the medical arena. It is important to note that the consent challenge is applicable to all patient information in all types of formats.

Provided that modern health information systems deal with patient data in electronic format, the consent challenge mutates into an electronic consent challenge. Patient consent has to be integrated at an application-native level in order for it to be useful. O’Keefe et al. [6] coined the term eConsent and listed the following as key attributes of an electronic patient consent management system:

- There is no single right solution to the eConsent challenge as an eConsent system needs to be tailored for the needs of different environments. For example, the degree to which eConsent plays the role of a gatekeeper, active audit or passive record should be customizable.
- The introduction of an eConsent system will formalize and potentially affect existing trust relationships between consumers and providers.
- A simple yet flexible model for eConsent, supported by education and training, is key to the acceptance and adoption of eConsent.
- A good set of default policies is key to ensuring that eConsent has a minimal impact on clinical workflow. For example, default policies for GP Clinics, Mental Health Clinics, Sexual Health Clinics and Hospitals could be established.
- The manner in which eConsent is integrated with existing systems can significantly affect its effectiveness.

The Semantic Challenge: Although the above mentioned features promise a solid foundation, they lack some of the practical considerations that may arise from a multi-vendor, multi-party health care information management infrastructure such as ours in Canada. When different parties and their

corresponding health information systems are required to collaborate, the information exchange mechanism may not be optimal and may require a lowest common denominator approach to information exchange. The major hurdle is not in the syntactic side, but rather in the semantic interoperability of information exchange. This impedance mismatch (between different vendor systems) poses one of the major challenges to universal patient consent management. To apply the governing rules of patient consent, one must have a sound semantic understanding of the consent as well as the data. We term this as the “Semantic Challenge”.

The semantic web technologies offer some primitive building blocks upon which a foundation for semantic consent management system can be laid. The traditional information exchange protocols utilize a static predefined schema in order to define the data structure of the shared information. All exchanged information is then governed (described) using the static schema. The semantic web offers a different approach to solve the information exchange problem between two dissimilar parties. Each dataset is accompanied by a corresponding vocabulary. Each vocabulary is described in a predefined and well understood format, therefore making it possible to describe data in ways that make its semantics explicit and hence discoverable automatically in software.

RDF is a standard model for data interchange on the Web. RDF has features that facilitate data merging even if the underlying schemas differ, and it specifically supports the evolution of schemas over time without requiring all the data consumers to be changed [7]. These properties of RDF make it an ideal candidate for management and exchange of medical information where the datasets (i) have a particular large set of consumers (ii) enjoy a long life span and (iii) are constantly evolving.

B. N3 Notation

Notation 3 (N3) is based on RDF standards and represents an analogous syntax to RDF/XML to represent data. Furthermore, N3 adds extra features (like rules and formulae) in order to process data and make inferences from facts in the data. N3 is designed as an alternative to RDF’s XML syntax with emphasis on (a) designed for a human-readability and (b) scribble language easy to work with (less verbose compared to RDF/XML) [8]. In N3, information is represented as a set of statements, where each statement is composed of a subject, verb and an object. Each statement ends with a period. For example an N3 statement:

`:thermo :temp :high.`

would translate into thermostat temperature is high. Similarly,

`:heating :power 0.`

would translate into heating power is zero. An N3 rule can be described by combining multiple N3 statements as such:

Tool	Description
cwm	Developed by the Tim Berners-Lee [4], cwm is a forward chain reasoner written in python for N3 and N3Logic. It can be used for querying, checking, transforming and filtering information. Furthermore, cwm is able to reason using a first order logic but without classical negation.
Euler	Euler is an inference engine that supports logic based proofs . It differs from cwm as it is a backward-chaining reasoner enhanced with Euler path detection mechanism [5].
Pychinko	Pychinko is an efficient Python implementation of the classic Rete pattern matching algorithm. It utilizes an optimized implementation of the algorithm to handle facts, expressed as triples, and process them using a set of N3 rules [4].

TABLE I
RDF/N3 REASONING TOOLS

`:thermo :temp :high. => :heating :power 0.`

means that if the thermostat temperature is high, then the heating power is zero. As stated before, an N3 formula adds a considerable amount of deduction and inference capabilities to a simple dataset. “By having rules and data in the same languages, N3 logic provides simplicity in syntax and completeness as rules can operate on themselves and anything written in N3 can be queried in N3” [4].

C. Euler

A reasoning engine is required in order to make use of the N3 rules. There are various RDF/N3 reasoners available. Table I highlights some of the most popular ones. Our decision to use Euler was motivated by the fact that it provides high level of integration of the core engine with high level programming languages used to build enterprise systems. Current Euler implementation supports Java, C#, Python, Javascript and Prolog. We used the Euler’s Java implementation for our application to integrate the reasoning engine into our demo application. The highly reusable design of Euler made the integration process painless. Furthermore, the Euler yap engine provides industrial level performance. Also to deal with large amount of data, Euler can translate data (triples) into SQL. Given the fact that Euler uses N3 notation, it makes it compatible with other reasoning engines based on N3 (such as W3C cwm). For the purpose of our project, the main motivational factors were: (a) support of N3 notation and rules (b) generation of the proof. The Euler engine satisfied both these conditions making it suitable for our project.

III. RELATED WORK

There are numerous studies dealing with electronic consent. However, these studies ignore the semantic aspect of information and focus mainly on security aspects [9], [10], [11]. O’Keefe et al. [6] undertake a feasibility study of electronic consent management system in medical arena. They expose various challenges faced by different consumer groups of electronic consent management systems. The study provides a sound set of recommendations for a generic implementation of a patient consent management system. Song et al. [12] introduce the notion of an e-consent object, encompassing

all relevant information concerning patient consent in the e-consent object. Lack of semantics is the biggest drawback of this model. The rules of consent are not expressed in any formal language and therefore are ambiguous at interpretation time.

Win et al. [13] describe an interface based approach through which patient consent can be expressed. The solution lacks organic growth as it hard codes the information and lacks the required flexibility for the user. Pruski et al. [14] propose e-CRL language designed with the following two goals in mind (a) facilitate capturing of patient consent information (b) formalize the expression of patient consent information. “The language has a well defined BNF based syntax and semantics defined based on first-order logic and set theory which allow eHealth systems to fully control the access to critical health data” [14]. Although the e-CRL language provides support for semantics, it lacks some important features such as proof generation. Furthermore, the defined language is not compatible with the RDF based solutions and approaches, making integration difficult.

IV. CONSENT POLICIES

There are different forms of consent that a patient may choose from. Coiera and Clarke [15] define four general forms of patient consent as follows:

- *General consent*: This is also generally known as “opt-in” which means that the patient agrees that any health care professional may access all of their health data for the purpose of providing care to them.
- *General consent with specific denial(s)*: This means that the patient specifically defines certain exceptions to their general consent policy. These exceptions may be related to particular data or to particular people.
- *General denial with specific consent(s)*: In this case, the patient generally denies access to their data except in specific circumstances. These may include disclosure for specific purposes or to specific people or the disclosure of certain types of information.
- *General denial*: This is the strictest level of consent a patient may have. In this case, the patient denies the use of their information for any future event irrelevant of the circumstances that may occur. This consent policy is also generally known as “opt-out”.




Fig. 1. System Overview

Pruski [14] use these four forms of consent while designing their rule-based language for describing patient consent policies. Additionally, Pruski outlines the key requirements that must be taken into consideration while dealing with patient consent. These are expressing who can access the data, the kind and the sensitivity of the data, the period for which a consent policy is valid, the purpose of why the data is being accessed, and what kind of consent is given to the user requesting the data.

In our work, we consider both the four forms of patient consent as well as some of the requirements proposed by Pruski. In our current prototype, we support five types of patient consent policies as follows:

- *opt-in*: As described above, a patient who has an opt-in policy allows any treating doctor to access their information. In our system, we also add other conditions that need to be met, such as that the health care professional should be a member of the admitting organization etc. These details will be discussed in Section V.
- *opt-in except for sensitive documents*: In this case, the patient allows access to all their information except for documents that are classified as sensitive. For example, these may include HIV test results or Sexually Transmitted Disease records.
- *opt-in except for certain people*: In this case, the patient allows access to all their information, but specifically denies certain individuals from accessing their data. We currently support exceptions to health care professionals only.
- *opt-out*: A patient with an opt-out policy explicitly denies access to all their information regardless of who is trying to access the data or why they are trying to access it.
- *opt-out with emergency override*: In this case, the patient agrees to grant access to their information only if it is an emergency situation.

V. CONSENTIR

This section describes the prototype of our proposed electronic patient consent management system, *Consentir*.

A. System Overview

Currently, *Consentir* assumes that all the information needed is already captured. That is, all information about the patient, their documents as well as their privacy policies. Similarly, we also assume that all the information about the different hospitals, doctors, and nurses is already present. In this sense, we focus on answering the question of whether a doctor or nurse can access a certain document. This is done by loading the facts available and our rule set into Euler, and then querying it to see if a certain doctor or nurse can access a certain document. Figure 1 shows an overview of how *Consentir* currently works.

If access is granted, then we present the user with the proof provided by Euler. If no proof is found, we check if we can find an explicit deny rule for the same actor and document in the original query. If so, we present the user with the proof of why access has been denied. Figure 2 shows our prototype tool. In this scenario, we are checking if Dr Smith has access to XRay1. In the “Figure” tab, a visual representation of all the facts related to this scenario is presented. These facts are extracted from the facts.n3 file described in the next section. Once the user clicks on “Check Access”, our tool would display Euler’s generated proof as shown in Figure 3. In this case, access has been granted as shown by the note in the top right corner. The proof shows the evidence used to reach this conclusion. It shows the different facts that have been used to satisfy the appropriate rule to either grant or deny access.

B. Facts in N3 Notation

To test our access policies, we chose a set of facts that can reflect all the situations we wish to test. Figure 4 shows all the facts we use to test our rules. There are four different types of actors: doctors, nurses, hospitals and patients. All hospitals have a “members only” policy where only members of the hospital are authorized to access patient information. Some hospitals enforce a “on shift” only policy where only members that are currently on shift may be able to access any patient information. Doctors and nurses can be members of more than one hospital, but can only be on shift in one hospital at a

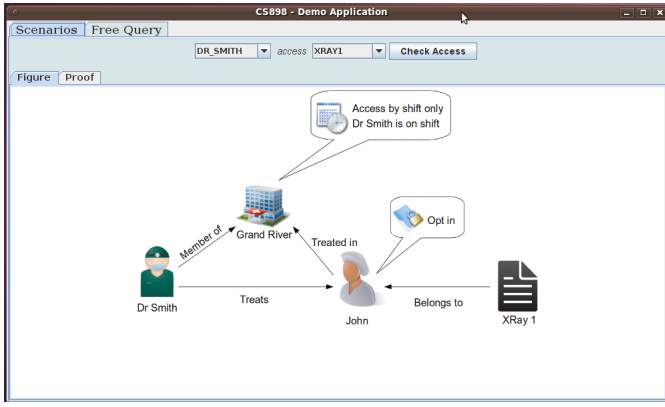


Fig. 2. Example Scenario

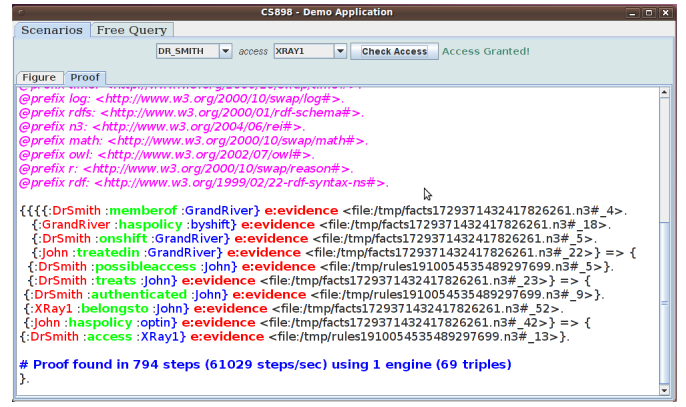


Fig. 3. Query Result & Proof of the Scenario in Figure 2

time. Patients are treated in a hospital, and own documents that reflect their medical information. Some documents are classified as sensitive. Currently, we only classify documents as sensitive or non-sensitive (the default). In future, we may wish to add different types of classifications such that a patient may have more flexibility in specifying which category of documents they grant access to. Additionally, a patient may be in an emergency situation. The default is that the patient is in a regular episode of care. As indicated in Section IV, we currently support five different types of patient consent policies. These are coded as optin, optinsens, optinexcep, and optout respectively. We assigned the different policies to the patients such that all of our designed access policies can be exhaustively tested. Appendix A shows the facts.n3 file we use to represent the facts shown in Figure 4.

C. Access Policies in N3 Notation

In our current prototype, we have designed the basic set of rules that can later be expanded for more sophisticated rule nesting. The complete rule file can be found in Appendix B. There are six possible rule conclusions in our rule set. The two main conclusions are allowing or denying access. However, we need the other rules to reach these conclusions. The following is the list of these six conclusions:

- **possibleaccess**: A person a has possible access to patient p if they are a member of the organization p is treated in. Additionally, if this organization has a by shift policy, then a must be on shift to have possible access.
- **authenticated**: A person a is authenticated to access a patient p 's information if they have possible access to p , and they are treating p . The reason possible access is separated from authentication is that in emergency situations, a non treating doctor might need to access this patient's information. In that case, they still need to satisfy the possible access rule.
- **access**: A person a is granted access to a document d if one of the following case is true:
 - a is authenticated to access p 's information, d belongs to p , and p has an opt in policy.

- a has possible access to p 's information, d belongs to p , p has an opt out policy with emergency override, and p is in an emergency situation.
- a is authenticated to access p 's information, d belongs to p , p has an opt in policy with exceptions, and a is not one of those exceptions.
- a is authenticated to access p 's information, d belongs to p , p has an opt in policy except for sensitive data, and d is not classified as sensitive.

- **cannotaccess**: This is the converse of the “possibleaccess” rule. A person a does not have possible access to p 's information if p is treated in o , but a is not a member of o . Similarly, if they are a member, but are not on shift in an organization that has a by shift policy, then a does not have possible access to p 's information.
- **notauthenticated**: This is the opposite of the “authenticated” rule. If a person a has possible access to p 's information, but a does not treat p , then a is not authenticated to access p 's information. If a cannot access p 's information in the first place according to the previous rule, then a is automatically not authenticated to access p 's information.
- **deny**: A person a is denied access to a document d in one of the following situations:
 - a is authenticated to access p 's information, d belongs to p , but p has an opt out policy.
 - d belongs to p , but a is not authenticated to access p 's information.
 - a has possible access to p , d belongs to p , p has an opt out policy with emergency override, and the situation is not an emergency.
 - a is authenticated to access p 's information, d belongs to p , p has an opt in policy except for sensitive information, and d is classified as sensitive.
 - a is authenticated to access p 's information, d belongs to p , p has an opt in policy with specific denials, and p has explicitly denied a access.

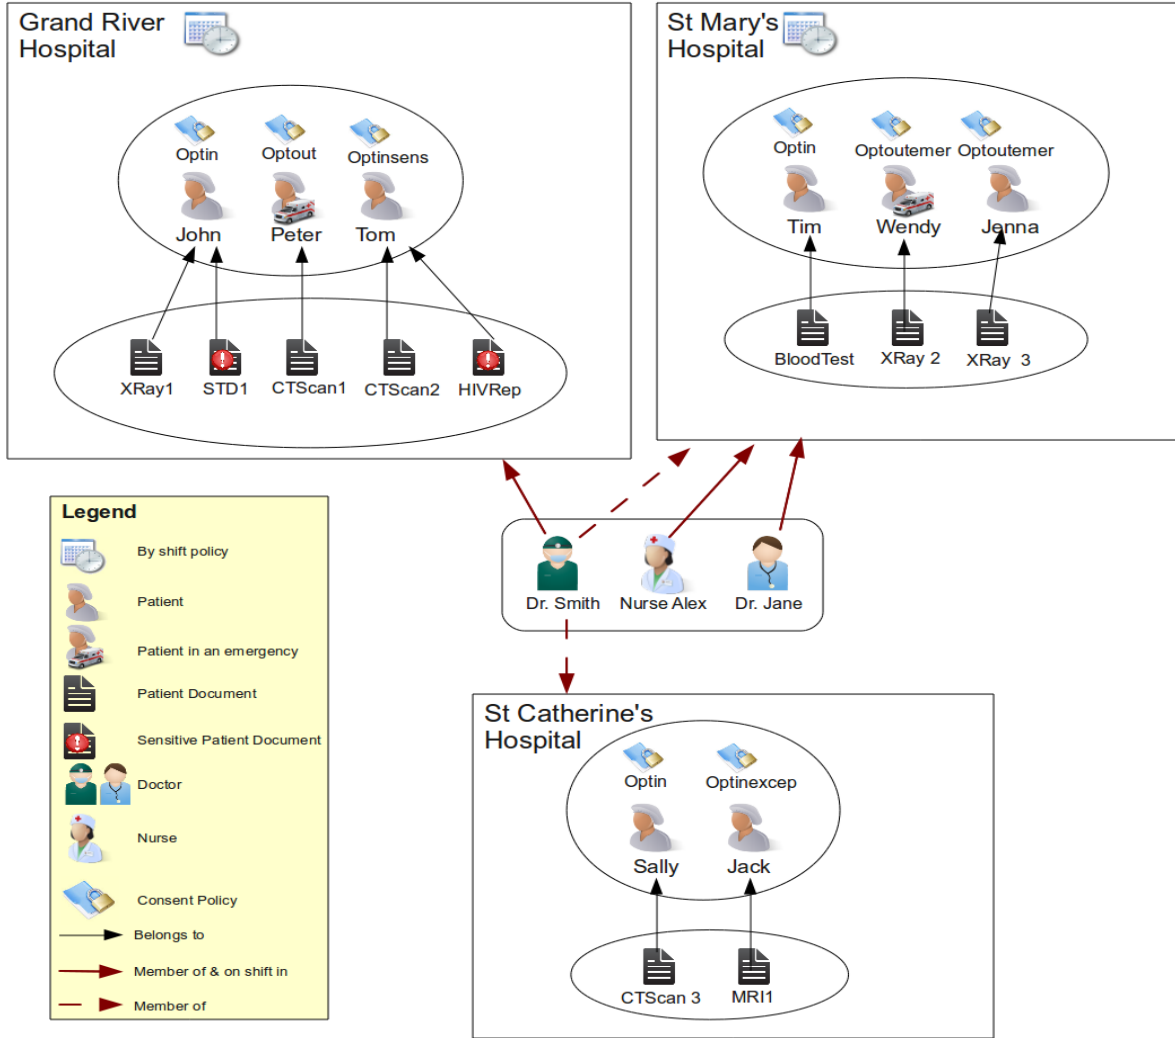


Fig. 4. Visual Representation of the facts.n3 File

D. Tested Scenarios

Given the above facts, we designed twelve different scenarios to test our different access policies. Table II summarizes these scenarios and their expected results. In this table, we use plain English to describe the query and facts instead of n3 notation for better illustration. These facts are based on the facts.n3 file shown in Appendix A, and represented in Figure 4. We write the queries in the form of “some person \rightarrow some document” which means that this person is trying to access this document, and we are trying to determine whether we should grant them access or not. As shown in Table II, we chose scenarios that show how the different consent policies behave in different situations. For example, Tom has an opt in policy with the exception of sensitive documents. When Dr Smith tries to access a non-sensitive document, CTScan2, he was granted access. However, when he tries to access a sensitive document, STD1, he was denied access. Similarly, we tested different combination of facts for the other policies to demonstrate how they produce different results in different

situations.

VI. DISCUSSION AND FUTURE WORK

Our current work is a simple prototype demonstrating the applicability of using ideas from the semantic web to the problem of reasoning with patient consent. Our prototype shows that using a reasoning engine such as Euler successfully provides the needed results. However, there are still some limitations to our solution. Some of them can be easily implemented as future work, and some of them are due to the nature of N3 notation, and the way Euler works.

Currently, we can only answer queries which have one answer. For example, we cannot answer queries in the form of “Who has access to a specific document?” This is because Euler returns the first answer it finds which satisfies the query. It does not look for all possible solutions. There may be ways to get around this which we are not currently aware of. As future work, it may be nice to provide such a querying functionality in the system.

Query	Relevant Facts	Expected Results
Dr Smith → XRay1	<ul style="list-style-type: none"> Dr Smith is a member of Grand River Grand River has a by shift policy Dr Smith is on shift in Grand River XRay1 belongs to John John is treated in Grand River Dr Smith treats John John has an opt in policy 	Dr Smith is granted access
Dr Smith → BloodTest	<ul style="list-style-type: none"> Dr Smith is a member of St Mary's St Mary's has a by shift policy BloodTest belongs to Tim Tim is treated in St Mary's Dr Smith treats Tim Tim has an opt in policy 	Dr Smith is denied access because St Mary's has a by shift policy and he is not currently on shift
Dr Smith → CTScan3	<ul style="list-style-type: none"> Dr Smith is a member of St Catherine's St Catherine's has a members only policy CTScan3 belongs to Sally Dr Smith treats Sally Sally has an opt in policy 	Dr Smith is granted access even though he is not on shift because St Catherine's does not have a by shift only policy
Dr Jane → BloodTest	<ul style="list-style-type: none"> Dr Jane is a member of St Mary's St Mary's has a members only policy Dr Jane is on shift at St Mary's BloodTest belongs to Tim Tim is treated in St Mary's Dr Smith treats Tim Tim has an opt in policy 	Dr Jane is denied access because she is not treating Tim
Dr Smith → CTScan1	<ul style="list-style-type: none"> Dr Smith is a member of Grand River Grand River has a by shift policy Dr Smith is on shift in Grand River CTScan1 belongs to Peter Peter is treated in Grand River Dr Smith treats Peter Peter has an opt out policy 	Dr Smith is denied access because Peter has an opt out policy
Dr Jane → XRay2	<ul style="list-style-type: none"> Dr Jane is a member of St Mary's St Mary's has a by shift policy Dr Jane is on shift in St Mary's XRay2 belongs to Wendy Wendy is treated in St Mary's Dr Jane treats Wendy Wendy has an opt out with emergency override policy Wendy is in an emergency situation 	Dr Jane is granted access since Wendy is in an emergency
NurseAlex → XRay2	<ul style="list-style-type: none"> Nurse Alex is a member of St Mary's St Mary's has a by shift policy Nurse Alex is on shift in St Mary's XRay2 belongs to Wendy Wendy is treated in St Mary's Wendy has an opt out policy with emergency override Wendy is in an emergency situation 	Nurse Alex is granted access even though she is not treating Wendy because Wendy is in an emergency
Dr Jane → XRay3	<ul style="list-style-type: none"> Dr Jane is a member of St Mary's St Mary's has a by shift policy Dr Jane is on shift in St Mary's XRay3 belongs to Jenna Dr Jane treats Jenna Jenna has an opt out policy with emergency override 	Dr Jane is denied access because Jenna has an opt out policy with emergency override, and Jenna is not in an emergency
Dr Smith → CTScan2	<ul style="list-style-type: none"> Dr Smith is a member of Grand River Grand River has a by shift policy Dr Smith is on shift at Grand River CTScan2 belongs to Tom Tom is treated in Grand River Dr Smith treats Tom Tom has an opt in policy with the exception of sensitive documents 	Dr Smith is granted access because CTScan2 is not classified as sensitive
Dr Smith → HIVRep1	<ul style="list-style-type: none"> Dr Smith is a member of Grand River Grand River has a by shift policy Dr Smith is on shift at Grand River HIVRep1 belongs to Tom Tom is treated in Grand River Dr Smith treats Tom Tom has an opt in policy with the exception of sensitive documents HIVRep1 is classified as sensitive 	Dr Smith is denied access because Tom has denied access to sensitive documents, and HIVRep1 is classified as sensitive
Dr Smith → STD1	<ul style="list-style-type: none"> Dr Smith is a member of Grand River Grand River has a by shift policy Dr Smith is on shift at Grand River STD1 belongs to John John is treated in Grand River Dr Smith treats John John has an opt in policy STD1 is classified as sensitive 	Dr Smith is granted access even though STD1 is sensitive because John has a full opt in policy
Dr Smith → MRI1	<ul style="list-style-type: none"> Dr Smith is a member of St Catherine's St Catherine's has a members only policy MRI1 belongs to Jack Jack is treated in St Catherine's Jack has an opt in policy with specific denials to certain people Jack has denied access to Dr Smith 	Dr Smith is denied access since Jack has explicitly denied him in his policy

TABLE II
SUMMARY OF THE TESTED SCENARIOS

To be able to address all medical situations, our rule set must be expanded. In this prototype, we demonstrated the basic consent policies. However, there are many variations that can be added. For example, allowing family members to ask for access. The basic rule construction would be the same, but we just need to build a rule ontology that accounts for all the possible situations. This leads to mentioning the need for more sophisticated rule writing in N3. N3 notation supports rule nesting which would support more elaborate policies. So far, we have focused on trying the basic features of writing rules in N3. The next step would be to expand the rule set to include more consent policies using rule nesting.

In our current prototype, we do not discover policy conflicts. This would be a very useful feature to add. First, it would allow us to discover any mistakes made while creating the rules. It would also be useful to detect conflicts between different levels of policies. For example, a hospital may have specific policies that allow access to patient information in a specific situation, while the province policies deny access. Detecting such conflicts and resolving them is important to achieve a fully functional electronic consent system. Similarly, implementing propagation scenarios where access can be propagated to another actor in certain situations would also be interesting. Additionally, we currently only support “access” or “deny” results. In order to be more specific, we plan to support more granular access levels in the future. For example, read, write, modify, delete etc.

Another challenge in using semantic web technologies in the medical domain is the open world versus closed world assumption. Currently, we do sort of a combination of both worlds. That is, if Euler cannot find a proof that grants access, we try to look for an explicit deny rule. However, explicitly defining a deny rule for every possible situation will result in an explosion in the number of rules. In the sensitive area of patient privacy, balancing those aspects is very important. Finally, instead of hard coding patient facts in N3 notation, having the relevant facts be automatically converted into N3 from an actual hospital database system would be more practical. That way, already existing information can be used and integrated within an electronic patient consent management system.

VII. CONCLUSION

In this paper, we presented *Consentir*, an electronic patient consent management system prototype. Our prototype is based on N3 notation, and the Euler reasoning engine. We demonstrated five different patient consent policies (opt in, opt in with exceptions to sensitive information, opt in with exceptions to certain people, opt out, and opt out with emergency override), and presented different scenarios to show how they affect access control in each situation. We tested twelve different scenarios which present different fact combinations to show how the different consent policies behave in different situations. Our prototype shows that using a semantic reasoning engine such as Euler works well to address the challenges of patient information access. The advantage of

Euler is that it can semantically reason with the data, and additionally, can generate a proof showing the evidence it found to support its conclusions. Such a proof can later be used as a verification token to be exchanged between different systems. We hope to expand this prototype by supporting more consent policies, and more elaborate rule nesting in N3 to achieve better support for patient consent management.

ACKNOWLEDGMENT

We would like to thank Dr. Helen Chen for her guidance throughout the development of this prototype.

REFERENCES

- [1] E. Kluge, “Informed consent and the security of the electronic health record (EHR): some policy considerations,” *International Journal of Medical Informatics*, vol. 73, no. 3, pp. 229–234, 2004.
- [2] “Notation 3 (n3): A readable RDF syntax,” Website, <http://www.w3.org/TeamSubmission/n3/>.
- [3] “Euler proof mechanism,” Website, <http://eulersharp.sourceforge.net/>.
- [4] T. Berners-lee, D. Connolly, L. Kagal, Y. Scharf, and J. Hendler, “N3logic: A logical framework for the world wide web,” *Theory Pract. Log. Program.*, vol. 8, no. 3, pp. 249–269, 2008.
- [5] G. Naudts, “An inference engine for rdf,” Master’s thesis, Open University of the Netherlands, 2003.
- [6] C. O’Keefe, A. Goodchild, P. Greenfield, A. Waugh, E. Cheung, and D. Austin, “Implementation of electronic consent mechanisms,” *Final Analysis Paper*, 2002.
- [7] “Resource description framework (rdf),” Website, <http://www.w3.org/RDF/>.
- [8] “Notation 3 – ideas about web architecture,” Website, <http://www.w3.org/DesignIssues/Notation3.html>.
- [9] J. Reid, I. Cheong, M. Henricksen, and J. Smith, “A novel use of rbac to protect privacy in distributed health care information systems,” in *ACISP’03: Proceedings of the 8th Australasian conference on Information security and privacy*. Berlin, Heidelberg: Springer-Verlag, 2003, pp. 403–415.
- [10] X. Chen, D. Berry, and W. Grimson, “Identity management to support access control in e-health systems,” in *4th European Conference of the International Federation for Medical and Biological Engineering*. Springer, 2009, pp. 880–886.
- [11] B. Blobel, “Authorisation and access control for electronic health record systems,” *International Journal of Medical Informatics*, vol. 73, no. 3, pp. 251–257, 2004.
- [12] H. Song, T. Win, and P. Croll, “Patient e-consent mechanism: Models and technologies,” *Proceedings of COLLECTeR, Melbourne, Australia*, 2002.
- [13] K. Win, H. Song, P. Croll, and J. Cooper, “Implementing patients consent in electronic health record systems,” *Proceedings of COLLECTeR, Melbourne, Australia*, 2002.
- [14] C. Pruski, “e-crl: A rule-based language for expressing patient electronic consent,” *eHealth, Telemedicine, and Social Medicine, International Conference on*, vol. 0, pp. 141–146, 2010.
- [15] E. Coiera and R. Clarke, “e-Consent: The design and implementation of consumer consent mechanisms in an electronic environment,” *Journal of the American Medical Informatics Association*, vol. 11, no. 2, p. 129, 2004.

APPENDIX A FACTS FILE

@prefix : < # >.

#Nurses and Doctors

:DrSmith :memberof :GrandRiver.
:DrSmith :onshift :GrandRiver.
:DrSmith :memberof :StMarys.
:DrSmith :memberof :StCatherines.

:DrJane :memberof :StMarys.
:DrJane :onshift :StMarys.

:NurseMary :memberof :GrandRiver.
:NurseAlex :memberof :StMarys.
:NurseAlex :onshift :StMarys.

#Hospital policies

:StMarys :haspolicy :byshift.
:GrandRiver :haspolicy :byshift.
:StCatherines :haspolicy :members.

#Patients

:John :treatedin :GrandRiver.
:DrSmith :treats :John.
:NurseMary :treats :John.
:Tim :treatedin :StMarys.
:DrSmith :treats :Tim.
:Peter :treatedin :GrandRiver.
:DrSmith :treats :Peter.
:NurseAlex :treats :John.
:Wendy :treatedin :StMarys.
:DrJane :treats :Wendy.
:Tom :treatedin :GrandRiver.
:DrSmith :treats :Tom.
:Jenna :treatedin :StMarys.
:DrJane :treats :Jenna.
:Sally :treatedin :StCatherines.
:DrSmith :treats :Sally.
:Jack :treatedin :StCatherines.
:DrSmith :treats :Jack.

#Patients' consent policies

:John :haspolicy :optin.
:Tim :haspolicy :optin.
:Peter :haspolicy :optout.
:Wendy :haspolicy :optoutemer.
:Tom :haspolicy :optinsens.
:Jenna :haspolicy :optoutemer.
:Sally :haspolicy :optin.
:Jack :haspolicy :optinexcep.

#Patients' documents

:XRay1 :belongsto :John.
:STD1 :belongsto :John.

:BloodTest :belongsto :Tim.
 :CTScan1 :belongsto :Peter.
 :XRay2 :belongsto :Wendy.
 :CTScan2 :belongsto :Tom.
 :HIVRep1 :belongsto :Tom.
 :XRay3 :belongsto :Jenna.
 :CTScan3 :belongsto :Sally.
 :MRI1 :belongsto :Jack.

#Patients' Situations

:Peter :hassituation :emergency.
 :Wendy :hassituation :emergency.

#Documents' sensitivity

:STD1 :hasnature :sensitive.
 :HIVRep1 :hasnature :sensitive.

#Denial of access to specific people

:Jack :denyaccess :DrSmith.

APPENDIX B RULES FILE

#These rules checks if a has possible access to patient p's records. a must be a member of the same organization that p is #treated in. Depending on the organization policy, member may have to be on shift.

?a :memberof ?o. ?o :haspolicy :byshift. ?a :onshift ?o. ?p :treatedin ?o. => ?a :possibleaccess ?p.

?a :memberof ?o. ?o :haspolicy :members. ?p :treatedin ?o. => ?a :possibleaccess ?p.

#If a has possible access to patient p's records, and a treats p then a is authenticated to access p's records

?a :possibleaccess ?p. ?a :treats ?p. => ?a :authenticated ?p.

#If a has possible access to d, and p has a full optin policy then a can access d

?a :authenticated ?p. ?d :belongsto ?p. ?p :haspolicy :optin. => ?a :access ?d.

#If a has possible access to d, and p has a an optout policy with emergency override and situation s is an emergency #situation, then a can access d even if the dr is not the treating doctor

?a :possibleaccess ?p. ?d :belongsto ?p. ?p :haspolicy :optoutemer. ?p :hassituation :emergency => ?a :access ?d.

#If p has an optin policy with exceptions, d belongs to p, and a is not one of those exceptions, then a is granted access
 #to d.

?a :authenticated ?p. ?d :belongsto ?p. ?p :haspolicy :optinexcep. <facts.n3> log:notIncludes ?p :denyaccess ?a => ?a :access ?d.

#If a has possible access to d, p has a opt in policy except for sensitive data, and d is not sensitive, then a #can access d

?a :authenticated ?p. ?d :belongsto ?p. ?p :haspolicy :optinsens. <facts.n3> log:notIncludes ?d :hasnature :sensitive => ?a :access ?d.

#DENIAL RULES

#No possible access if a is not a member of the same organization p is treated in

?p :treatedin ?o. <facts.n3> log:notIncludes ?a :memberof ?o => ?a :cannotaccess ?p.

#No possible access if a is not on shift in an organization that has a by shift policy

?p :treatedin ?o. ?a :memberof ?o. ?o :haspolicy :byshift. <facts.n3> log:notIncludes ?a :onshift ?o => ?a :cannotaccess ?p.

#Not authenticated if a has possible access but does not treat p

?a :possibleaccess ?p. <facts.n3> log:notIncludes ?a :treats ?p => ?a :notauthenticated ?p.

#If a has possible access to d, and p has a an optout policy with emergency override and situation s is NOT an emergency situation, then a is denied access.

?a :possibleaccess ?p. ?d :belongsto ?p. ?p :haspolicy :optoutemer. <facts.n3> log:notIncludes ?p :hassituation :emergency => ?a :deny ?d.

#If p has an opt in policy with exceptions and explicitly denies a doctor then access to this document is denied

?a :authenticated ?p. ?d :belongsto ?p. ?p :haspolicy :optinexcep. ?p :denyaccess ?a => ?a :deny ?d.

If a is authenticated, and p has a full optout policy, and #document d belongs to patient p then a is denied access
#to d

?a :authenticated ?p. ?d :belongsto ?p. ?p :haspolicy :optout. => ?a :deny ?d.

#Not authenticated if a cannot access p (even though a treats p (maybe not on shift)).

?a :cannotaccess ?p. => ?a :notauthenticated ?p.

#If a is not authenticated to treat p, then access is denied

?a :notauthenticated ?p. ?d :belongsto ?p => ?a :deny ?d.

#If a is authenticated, d belongs to p, p has a an optin #except to sensitive data, and d is sensitive then a is denied
#access

?a :authenticated ?p. ?d :belongsto ?p. ?p :haspolicy :optinsens. ?d :hasnature :sensitive => ?a :deny ?d.